

FILE I/O

CS 23200

Big Picture

- ✓ Developing programs on *nix computers
- C Language
 - ✓ Familiar aspects of C (variables, operators, basic I/O, control flow, functions)
 - ✓ Pointers
 - ✓ Structures and related constructs
 - File operations
 - Standard library functions
 - Multi-file programs
- *nix tools

Outline

- Manipulating I/O from the shell
 - Redirection
 - Piping
- C functions for I/O from standard in/out
- C functions for general file I/O
 - Formatted I/O
 - Unformatted text I/O
 - Binary files
 - File positioning

Standard In, Out, Error

- Standard input
 - Typically the stream of keyboard input from the terminal
- Standard output
 - Programs often print information to the standard output stream (using printf, e.g.)
 - Typically displayed to the terminal
- Standard error
 - Programs print error messages to the standard error stream
 - Also typically displayed to the terminal

Input Redirection

- What if we want to automate standard input?
 - ▣ Putting theBomb passwords in a text file, e.g.
- Use the input redirection operator (<) in the shell

```
[shell prompt $] ./myProgram <inputFile.txt
```
- Will take the contents of inputFile.txt as if they were entered on the keyboard

Input Redirection

```
[shell $] ./myProgram <inputFile.txt
```

- Note that myProgram doesn't know the redirection has happened
 - ▣ <inputFile.txt is not a command line argument
 - ▣ The <inputFile.txt is processed by the shell before the program starts

Output Redirection

- Instead of displaying output on the screen, can redirect it to a file

```
[shell $] ./myProgram >consoleOutput.txt
```

- To redirect standard error as well:

```
[shell $] ./myProgram >consoleOutput.txt 2>&1
```

sends stdout to consoleOutput.txt

redirects
stderr (2) into
stdout (1)

- ">>" append the results to the file

Piping

- Several *nix tools do processing on standard input and write results to standard output
- Can use the output of one program as the input of another program using piping

```
[shell $] ./prog1 | ./prog2
```

- ▣ Standard output of prog1 is the standard input of prog2
- Example: checking your primes project

```
[shell $] ./primes 10 20 | ./checkPrimes 10 20
```

tee

- tee is a *nix program that lets you capture output to a file and print it to the screen

- Usage:

```
./myProg arg1 arg2 | tee consoleOutput.txt
```

- tee takes the standard input stream and copies it to a file and to the screen

Outline

- Manipulating I/O from the shell
 - ▣ Redirection
 - ▣ Piping
- **C functions for I/O from standard in/out**
- C functions for general file I/O
 - ▣ Formatted I/O
 - ▣ Unformatted text I/O
 - ▣ Binary files
 - ▣ File positioning

Single-character I/O

- **Input:** `int getchar(void);`
 - ▣ Returns the next unprocessed character from standard input
 - ▣ Returning int permits returning EOF, which indicates end-of-file

```
int c;  
c = getchar();  
if(c != EOF){  
    ...
```

Single-character I/O

- **Output:** `int putchar(int c);`
 - ▣ Writes c to standard output
 - ▣ Normally, returns c
 - ▣ On error, returns EOF

```
const char *str = "A string to write \n";  
const int n = strlen(str);  
int i;  
  
for(i=0; i<n; i++){  
    if(putchar(str[i]) == EOF){  
        /* do error handling */  
        ...  
    }  
}
```

Exercise: Single-character I/O

Exercise 7-1 (adapted) from K&R:

- Write a program that takes one command-line argument: "L" or "U"
 - ▣ If "L" the program will convert to lower case
 - ▣ If "U" the program will convert to upper case
- The program should read from standard input, doing the case conversion before echoing back to standard output
- Use "tolower(c)" and "toupper(c)"

Solution: Single-character I/O

```
#include <stdio.h>

int main(int argc, char** argv){
    int c, goLower;
    if(argc < 2){
        fprintf(stderr, "USAGE: %s [lower|upper]\n", argv[0]);
        return -1;
    }
    goLower = !strcmp("lower", argv[1]);
    c = getchar();
    while( c != EOF ){
        if (goLower)
            putchar(tolower(c));
        else
            putchar(toupper(c));
        c = getchar();
    }
    return 0;
}
```

Good ol' printf

- Use printf to print to standard output
- printf takes a format string and some arguments

```
printf("%s is a string.\n"
      "%d is an int.\n"
      "%f is a double.\n",
      "\"This\"", -78, 54.2);
```

↖ ↗ No commas; these string literals are concatenated by the compiler

Good ol' printf

- Format strings are often string literals

```
double x = -5.4;
printf("%.12f", x);
```

- But they can also be string variables

```
const char *dblFormat = "%.12f";
double x;
...
printf(dblFormat, x);
```

Advanced printf

- printf format specifiers are very flexible, allowing lots of output formats

the % sign; required → %[-][min][.prec][h or l]type

the character for the type (e.g., d, f, s); required

minimum field width; padded with spaces if necessary; adds spaces on the left unless...

... there is a negative sign here; then the spaces are padded on the right

h for short, l for long;
%hd is a short,
%ld is a long,
%lf is a long double

Advanced printf

- printf format specifiers are very flexible, allowing lots of output formats

the % sign; required → %[-][min][.prec][h or l]type

the character for the type (e.g., d, f, s); required

precision behavior depends on type:

- **string:** **maximum** number of characters to print
- **floating point:** **exact** number of digits after the decimal point
- **integer:** **minimum** number of digits (will add leading zeroes if necessary)

minimum field width; padded with spaces if necessary

Fixed-width Fields

- Fixed-width fields are a common goal

Jane	Doe	12	3.864
George	Smith	7349	278.100

not

Jane	Doe	12	3.864
George	Smith	7349	278.1

Fixed-width Fields

Jane	Doe	12	3.864
George	Smith	7349	278.100

- Potential problem 1: data is too narrow
 - Solution: use minimum width to pad with spaces

```
printf("%-8s %-8s %8d %8.3f\n", ...);
```

Why the negative signs?

Why the .3 for the floating point?

Fixed-width Fields

```
Jane      Doe      12      3.864
George    Smith    7349    278.100
Rumpelstiltskin Jones  1816473825 37857.357
Alexande Johnson      139 574839.320
```

□ Potential problem 2: data is too wide

▣ Solution depends on type

▣ String: use precision to specify maximum width

```
printf("%-8.8s %-8.8s %d %8.3f\n", ...);
```

```
Rumpelst Jones  1816473825 37857.357
Alexande Johnson      139 574839.320
```

Fixed-width Fields

```
printf("%-8.8s %-8.8s %d %8.3f\n", ...);
```

```
Rumpelst Jones  1816473825 37857.357
Alexande Johnson      139 574839.320
```

□ Potential problem 2: data is too wide

▣ No way to specify maximum width for numerical types

▣ Must set the field width based on the maximum length value you will be printing

▣ Not very robust

▣ Should check values before printf if formatting is crucial

```
if(num >= 100000000 || num <= -100000000){
    /* num is too wide for an 8-character field */
    /* print an error message */
    ...
}
```

printf exercise

```
struct scoreInfo{
    char* name;
    int id;
    double score;
};
...
const int n = 200;
struct scoreInfo scores[n];
...
/* scores is initialized */
...
```

□ Write code to print the data in scores

▣ Use one line for each item in the array

▣ Use alignment so the data for a given field always starts at the same offset

printf exercise: solution

```
struct scoreInfo{
    char* name;
    int id;
    double score;
};
...
const int n = 200;
struct scoreInfo scores[n];
...
/* scores is initialized */
...

int i;
for(i=0; i<n; i++){
    printf("%9d %-20.20s %8.3f\n",
           scores[i].id, scores[i].name,
           scores[i].score);
}
```

Formatted Input

- scanf is the input counterpart to printf
 - printf prints output to standard out
 - scanf grabs input from standard in

```
int intOne, intTwo;
double dblOne;

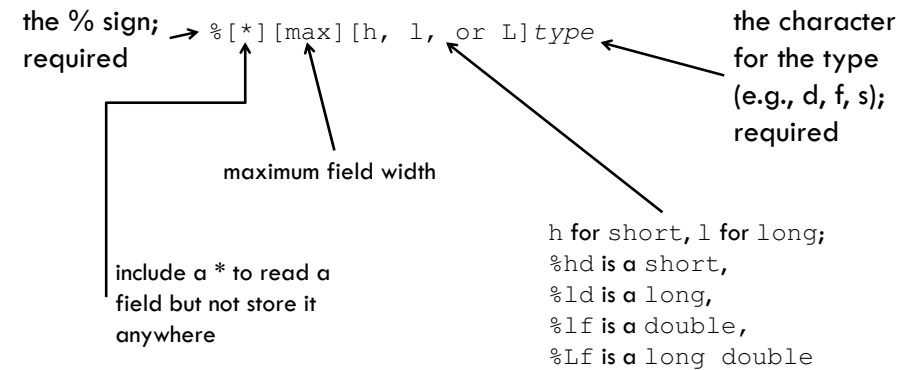
printf("Enter two integers and a floating point:\n");
scanf("%d %d %lf", &intOne, &intTwo, &dblOne);
```

use %lf for double,
%f for float

pass **addresses** for storing the input

scanf Format Specifiers

- scanf format specifiers are slightly different from printf's



scanf and Whitespace

- Blanks or tabs in format string tell scanf to gobble up as much white space (newlines, spaces, tabs, etc.) as it can

```
scanf("%d %d", ...);
/* same as */
scanf("%d %d", ...);
```

- scanf will gobble up white space before parsing a field

```
/* above examples are the same as */
scanf("%d%d", ...);
```

- Use %c specifier or getchar() to get whitespace characters

- scanf("%d %d", ...) will correctly process...

- 234 789
- 234 789
- 234 789

scanf

- %s reads in characters until a whitespace
- Ordinary characters (not whitespace or format specifiers) should match the input

```
char name[9];
scanf("name: %8s", name);
```

succeeds on:

name: John

fails on:

Name: John

scanf

- NEVER use %s in scanf without specifying a maximum width
 - ▣ This would allow the user to crash your program (or execute their own code) by specifying a long string
- scanf returns the number of successfully parsed fields
 - ▣ Stops processing input after the first failure
- Remember: scanf arguments should be pointers (i.e., addresses)

scanf exercise

- A user enters 3 integers per line for n lines
- Compute and print the column sums
 - ▣ e.g., the sum of the first integer on each line is the first column sum

scanf exercise solution

- Remember: pass ADDRESSES to scanf

```
const int numCols = 3;
int sums[3] = {0,0,0};
int lineNum, colNum;

for(lineNum=0; lineNum<n; lineNum++){
    int vals[numCols];
    scanf("%d %d %d%c", vals, vals+1, vals+2);
    for(colNum=0; colNum<numCols; colNum++){
        sums[colNum] += vals[colNum];
    }
}
printf("%d %d %d\n", sums[0], sums[1], sums[2]);
```

<http://stackoverflow.com/questions/3744776/simple-c-scanf-does-not-work>

Outline

- Manipulating I/O from the shell
 - ▣ Redirection
 - ▣ Piping
- C functions for I/O from standard in/out
- **C functions for general file I/O**
 - ▣ **Formatted I/O**
 - ▣ **Unformatted text I/O**
 - ▣ **Binary files**
 - ▣ **File positioning**

General File I/O

□ General pattern

▣ Open a file

- Returns a FILE*
- FILE* identifies the file in subsequent operations (read, write, close)
- "Opening" can include creating the file if it doesn't exist

▣ Read/write from the file

▣ Close the file

fopen

□ FILE* fopen(char *filename, char *mode);

▣ filename can be a relative or absolute path

▣ mode

- "r" for read
- "w" for write
- "a" for append
- "r+" for read and write
- BE CAREFUL: "w" will create a blank file with the given name, overwriting an existing file

▣ Returns NULL on error

- ALWAYS check for NULL
- More on error handling later

fclose

□ int fclose(FILE *stream);

▣ Closes the file indicated by stream

▣ Returns 0 upon success, EOF upon error

Reading and Writing

□ Different sets of functions for reading and writing

□ One set: fprintf and fscanf

- ▣ Work like printf and scanf
- ▣ One additional argument: first argument is the FILE*

□ printf and scanf are actually special cases of fprintf and fscanf

- ▣ Standard input, standard output, and standard error each have global FILE* variables
- ▣ `stdin`, `stdout`, `stderr`

fprintf Example

```
int main(int argc, char* argv[]){
    FILE *fp;
    if(argc < 2){
        fprintf(stderr, "Missing file name.\n");
        return -1;
    }
    fp = fopen(argv[1], "w");
    if(fp == NULL){
        fprintf(stderr, "Could not create file (%s).\n",
            argv[1]);
        return -2;
    }
    fprintf(fp, "This is a number: %d\n", 17);
    fprintf(fp, "Another number: %8.4f\n", 3.141592);
    fclose(fp);
    return 0;
}
```

fscanf Example

- Let "matrix.txt" be a file with the following format:

number of rows
number of columns
data, one line for each row

```
2
3
3.42 -47.0 9.3
0.0 8.38 -374.2
```

```
const char* filename = "matrix.txt";
int numRows, numCols, res, row, col;
double **matrix;
FILE *fp = fopen(filename, "r");
if(fp == NULL){
    fprintf(stderr, "Could not open file (%s).\n", filename);
    return -1;
}
res = fscanf(fp, "%d", &numRows);
if(res != 1){
    fprintf(stderr, "File format error.\n");
    return -2;
}
res = fscanf(fp, "%d", &numCols);
... /* check res and create matrix */
for(row=0; row<numRows; row++){
    for(col=0; col<numCols; col++){
        fscanf(fp, "%lf", &(matrix[row][col]));
    }
}
fclose(fp);
```

Exercises

- Two programs:
 - ▣ Write your birthday to a file "birthday.txt"
 - ▣ Read your birthday from the file you just wrote
 - Should have month, date, year variables

A Solution for Writing

```
int main(){
    const char *filename = "birthday.txt";
    FILE *fp = fopen(filename, "w");

    if(fp == NULL){
        fprintf(stderr, "Could not open file (%s)\n",
            filename);
        return -1;
    }

    fprintf(fp, "%d/%d/%d\n", 3, 28, 1928);
    fclose(fp);

    return 0;
}
```

A Solution for Reading

```
int main(){
    const char *filename = "birthday.txt";
    FILE *fp = fopen(filename, "r");
    int month, date, year, res;
    if(fp == NULL){
        fprintf(stderr, "Could not open file (%s)\n",
            filename);
        return -1;
    }
    res = fscanf(fp, "%d/%d/%d\n", &month, &date, &year);
    if(res != 3){
        fprintf(stderr, "File format error.\n");
        return -2;
    }
    fclose(fp);
    return 0;
}
```

Another Example

- Calculate the average of students' scores
- Scores are stored in a text file
- Format in the file:
student_name score
- get_average.c

Landscape of I/O Functions

	Formatted	Unformatted	Binary
Input	fscanf (scanf is a special case)	single char: fgetc whole line: fgets	fread
Output	fprintf (printf is a special case)	single char: fputc more: fputs	fwrite

Unformatted Text I/O

□ Single characters

- ▣ `int fgetc(FILE *stream);`
 - Returns next character from stream (EOF if end of stream)
- ▣ `int fputc(int c, FILE *stream);`
 - Writes `c` to stream
 - Returns the character written, or EOF for error

□ Peeking

- ▣ Might want to stop processing if the next character satisfies some condition
- ▣ `fgetc` to look at the character, then `ungetc` to put it back
- ▣ `int ungetc(int c, FILE *stream);`
- ▣ Only guaranteed for one character

Unformatted Text I/O

□ More than single characters

- ▣ `char* fgets(char *s, int n, FILE *stream);`
 - Reads characters into `s` until...
 - A `'\n'` is reached
 - or `n-1` characters have been read
 - `s` will always be null-terminated
 - `s` might contain a newline `'\n'`
 - Returns `s`, or NULL if end-of-file error occurs
- ▣ `int fputs(const char *s, FILE *stream);`
 - Writes `s` to stream
 - `s` need not contain a newline
 - Returns EOF for error

A Common File Input Pattern

□ Loop over lines while `fgets` does not return NULL

- ▣ Process the line that `fgets` just read
 - Often use `sscanf` to extract int's or double's

```
int sscanf(const char* source, const char* format, ...);
```

- Works like `fscanf`
- Processes the string `source` (instead of input from a file)

The Reasoning

- Why use `fgets` and `sscanf` instead of using `fscanf` directly?
 - ▣ With `fgets`, an erroneous line is isolated from other lines
 - ▣ The whitespace issue
 - Expecting 3 int's per line: `fscanf(fp, "%d %d %d\n", ...);`
 - If one line is missing an int, this will grab the first int from the next line as the third int
 - ▣ Can look through the line multiple times if needed
- Disadvantage: lose the automatic position advancement within the line
 - ▣ Will discuss how to deal with this later

Example

- Get users inputs for a list of integers
- Calculate the sum and print it
- Keep looping until user type the empty string
- Number of integers can be varying

- `calculate_sum.c`

An Exercise

- The file "data.txt" contains the following information:
 - ▣ First line: the number of rows that follow
 - ▣ Each subsequent line: a list of whitespace-delimited integers
 - Each list can be a different length
- Example: each line contains Twitter followers (integer ids) for one person
- Read in the data and store each row in its own array of int's (for later processing)

An Exercise

- First line: the number of rows that follow
- Other lines: lists of integers
 - Each list can be a different length
- Read in the data and store each row in an array of int's (for later processing)

- **The process for approaching any problem:**
 - ▣ What are the big steps?
 - Then break each of those down into smaller steps
 - ▣ What data structures will I need for those steps?
 - ▣ Often, the first answers to these questions will need to be revised as you work out the details

The Steps

- Open the file and verify the file pointer
- Read the number of lines
- Create an array with one element for each line
- Loop over the lines
 - ▣ Read in the line
 - ▣ Determine how many int's there are
 - ▣ Allocate space for those int's
 - ▣ Read the int's into that space
- Close the file

We don't currently have the tools to do these, but we can still write the other code

The Data Structure

□ What information will we need from the file?

```
struct arrayOfVectors{
    int numRows;

    /* an array of size numRows */
    int *rowSizes;

    /* an array of size numRows;
       theData[i] points to an array
       of int's of size rowSizes[i] */
    int **theData;
};
```

```
res = fscanf(fp, "%d", &pResult->numRows);
if(res != 1 || pResult->numRows > 1000000 ||
    pResult->numRows <= 0){
    fprintf(stderr, "Invalid number of lines.\n");
    fclose(fp);
    free(pResult);
    return NULL;
}
pResult->theData = malloc(sizeof(int*) * pResult->numRows);
if(pResult->theData == NULL){
    fprintf(stderr, "Could not allocate memory.\n");
    fclose(fp);
    free(pResult);
    return NULL;
}
pResult->rowSizes = malloc(sizeof(int) * pResult->numRows);
if(pResult->rowSizes == NULL) {
    fprintf(stderr, "Could not allocate memory.\n");
    free(pResult->theData);
    fclose(fp);
    free(pResult);
    return NULL;
}
```

```
struct arrayOfVectors* readData(const char* filename){
    struct arrayOfVectors* pResult = NULL;
    int lineIndex, res;
    const int bufSize = 10000;
    char line[bufSize];
    FILE *fp = NULL;

    pResult = malloc(sizeof(struct arrayOfVectors));
    if(pResult == NULL){
        fprintf(stderr, "Could not allocate memory.\n");
        return NULL;
    }

    fp = fopen(filename, "r");
    if(fp == NULL){
        fprintf(stderr, "Could not open the file (%s).\n",
            filename);
        free(pResult);
        return NULL;
    }
}
```

Noticing a Pattern?

- One line of processing
- Check for errors
 - ▣ Error handling
- Next line of processing
- Check for errors
 - ▣ Error handling
- ...

- While this can get tedious, it's the only way to ensure secure, robust code

```

struct arrayOfVectors* readData(const char* filename){
    struct arrayOfVectors* pResult = NULL;
    int lineIndex, res;
    const int bufSize = 10000;
    char line[bufSize];
    FILE *fp = NULL;
    ...
    for(lineIndex=0; lineIndex<pResult->numRows; lineIndex++){
        fgets(line, bufSize, fp);

        /* by using fgets, we have a string to pass
           to createIntArray, which can manipulate it
           as necessary (in contrast to using fscanf) */
        pResult->theData[lineIndex] =
            createIntArray(line, pResult->rowSizes + lineIndex);
    }

    fclose(fp);
    return pResult;
}

```

Landscape of I/O Functions

	Formatted	Unformatted	Binary
Input	fscanf (scanf is a special case)	single char: fgetc whole line: fgets	fread
Output	fprintf (printf is a special case)	single char: fputc more: fputs	fwrite

Binary File I/O

- Using binary file I/O will make your code system-dependent

- ▣ Might reduce file size

```

size_t fread(void *ptr, size_t size,
             size_t nobj, FILE *stream);
size_t fwrite(const void *ptr, size_t size,
             size_t nobj, FILE *stream);

```

- Returns number of objects read/written

File Positioning

- The functions so far start at the beginning of the file and read/write in order
- What if you need to go back or skip ahead?
- `int fseek(FILE *stream, long offset, int origin);`
 - ▣ Moves the file position
 - ▣ origin:
 - `SEEK_SET` : beginning of file
 - `SEEK_CUR` : current position
 - `SEEK_END` : end of file
 - ▣ Returns non-zero on error

Summary

- Manipulating I/O from the shell
 - ▣ Redirection
 - ▣ Piping
- C functions for I/O from standard in/out
- C functions for general file I/O
 - ▣ Formatted I/O
 - ▣ Unformatted text I/O
 - ▣ Binary files
 - ▣ File positioning

Choosing an Input Method

